

# RTOS Scheduling

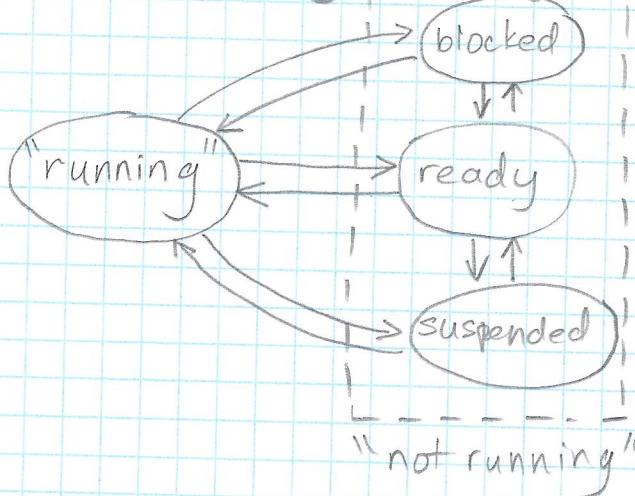
RJS-W3-1

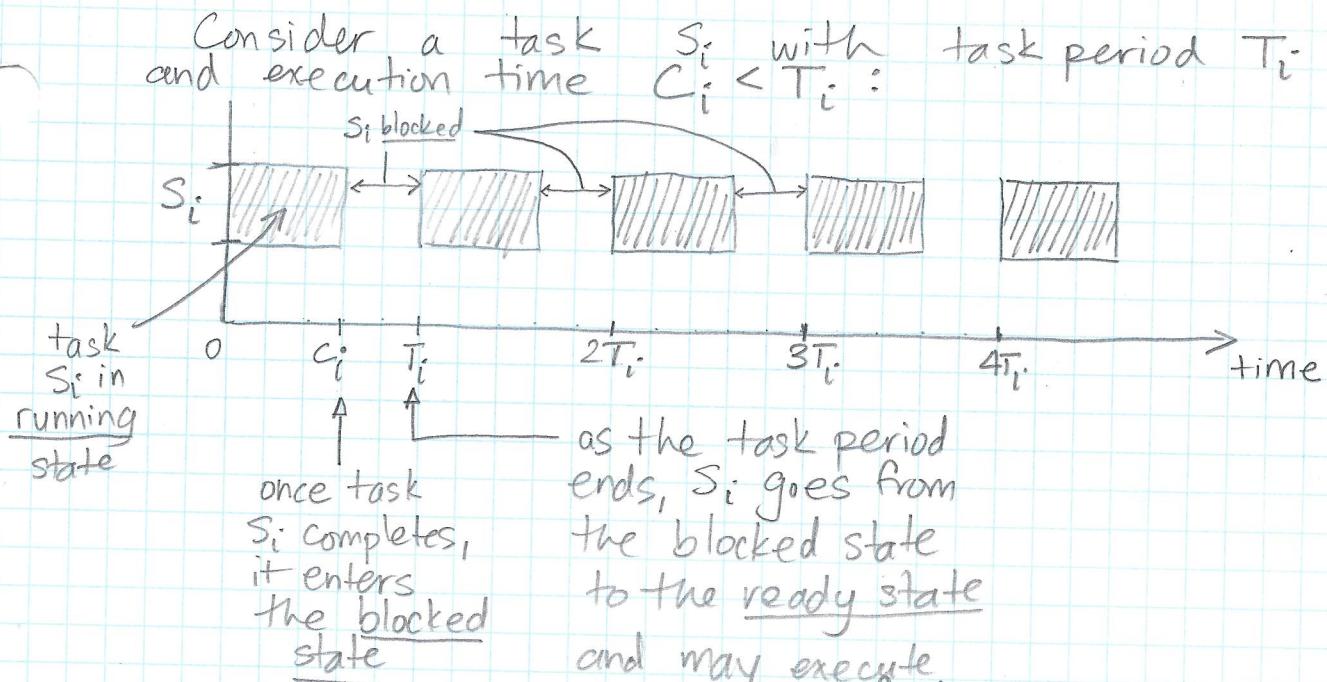
## Recap

- For our development here, we will use the following symbols frequently:
  - $C_i$  : the execution time needed to complete task/service  $S_i$ ; note that this does not include preemption delays — it is the "raw" execution time needed to complete  $S_i$  without interruptions.
  - $T_i$  : the task period: the fixed length of time between task calls or "releases".
  - $D_i$  : the deadline for task  $S_i$ ; we must have  $D_i \leq T_i$  in order to have no missed deadlines.

## Task States

- Generally tasks are either "running" or "not running" — if a task is not running, it may be "blocked", "suspended", or "ready".
- IMPORTANT: When the scheduler is selecting which task to run, it selects from amongst those tasks in the ready state or already running, based on priority.



EXAMPLE:

is the highest priority task in the ready state, it begins execution immediately

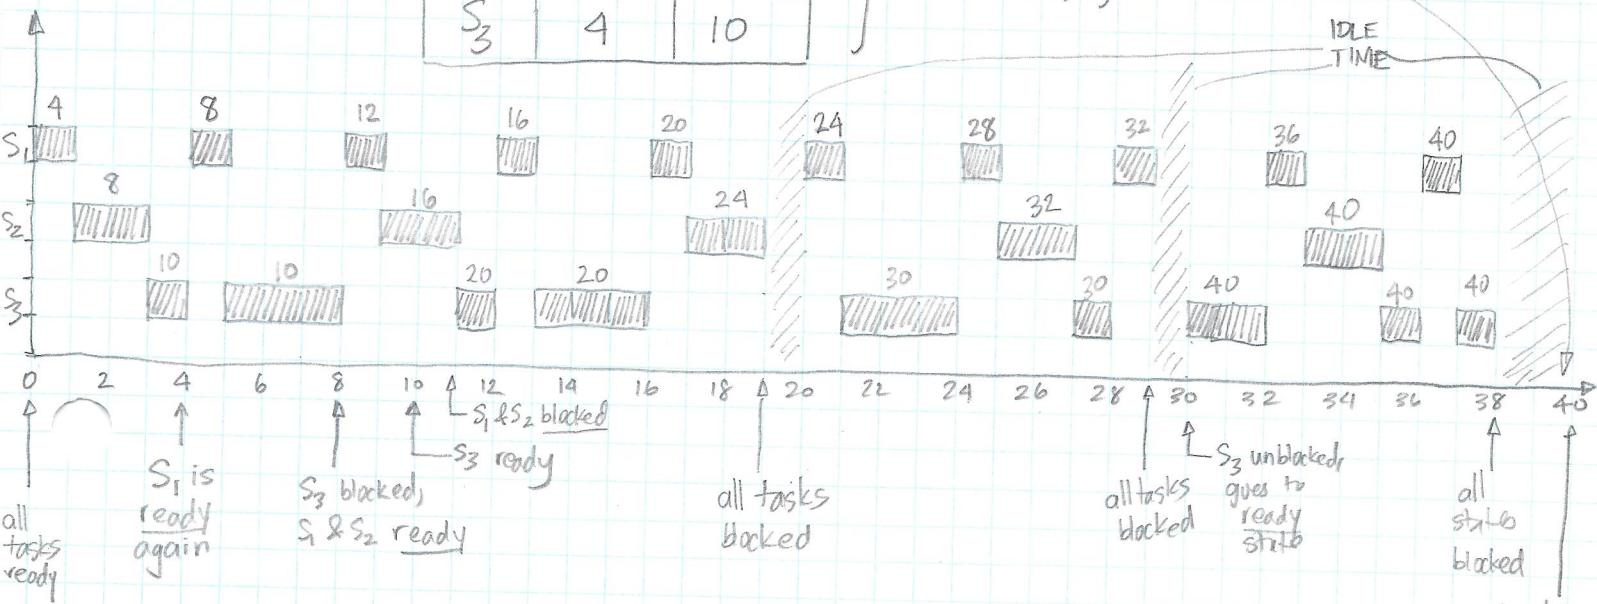
EXAMPLE:

Apply the RM fixed priority assignment policy to the following set of tasks:

	$C_i$	$T_i$
$S_1$	1	4
$S_2$	2	8
$S_3$	4	10

prior ( $S_1$ ) > prior ( $S_2$ ) > prior ( $S_3$ )

$$\text{LCM}(4, 8, 10) = 40$$



- in this example, we see the presence of idle time, suggesting that the CPU is not fully utilized,

- also note that the behaviour of this system is captured completely over the LCM time of all task periods — we can see that there are no missed deadlines

## Estimating Feasibility

- In RTOS theory, there is a sufficient (but not necessary) simple test of scheduling feasibility; that is, given  $N$  tasks  $S_i$ , can scheduling be achieved with no missed deadlines using the RM policy? This test is called the RM LUB (least upper bound):

- Theorem: Given the utility of the CPU,  $U := \sum_{i=1}^N \left( \frac{C_i}{T_i} \right)$ , if  $U \leq N(2^N - 1)$ ,

then RM priority assignment will produce no missed deadlines. ■

(stated w/o proof)

- in the previous example,

$$\begin{aligned} U &= \sum_{i=1}^3 \left( \frac{C_i}{T_i} \right) = \frac{C_1}{T_1} + \frac{C_2}{T_2} + \frac{C_3}{T_3} = \frac{1}{4} + \frac{2}{8} + \frac{4}{10} \\ &= \frac{1}{2} + \frac{2}{5} = \frac{5+4}{10} = \frac{9}{10} = 90\% \end{aligned}$$

and

$$N(2^N - 1) = 3(2^3 - 1) \approx 0.78$$

and so

$$U = 0.90 \neq 0.78$$

Showing that the RM LUB is conservative; because we have shown that RM policy, in fact, works with no missed deadlines.

- note that our example, however, had close to 100% CPU utilization (almost no idle time).
- if we "relax" our parameters further, say, set  $C_3 = 2$ , in order to give the scheduler "more room", then  $U$  becomes

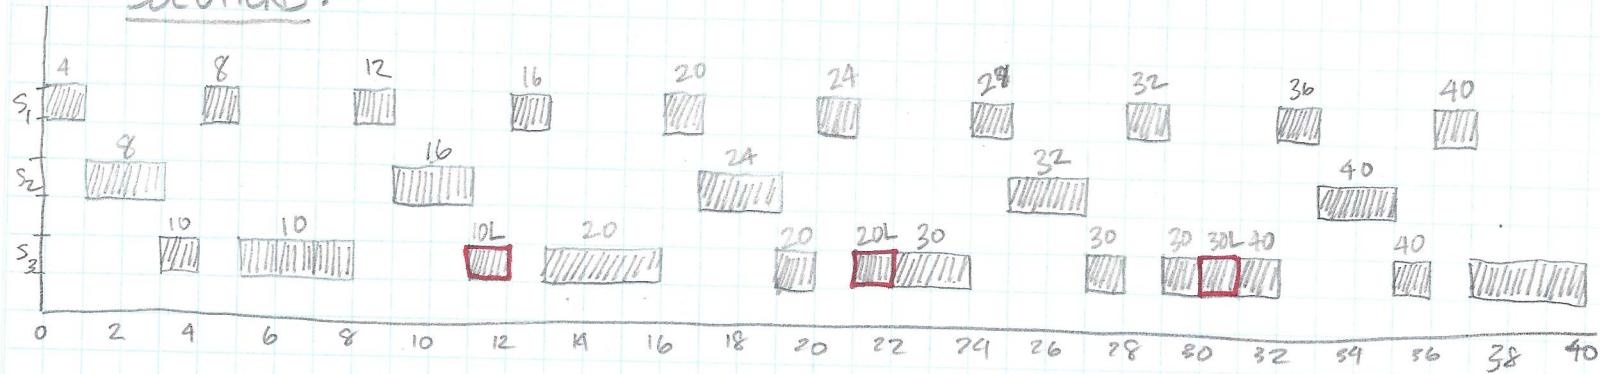
$$U = \frac{C_1}{T_1} + \frac{C_2}{T_2} + \frac{C_3}{T_3} = \frac{1}{4} + \frac{2}{8} + \frac{2}{10} \xrightarrow{\text{reduced from 4}} = \frac{1}{2} + \frac{2}{10} = \frac{5+2}{10} = \frac{7}{10} = 0.70.$$

in this case,  $U = 0.70 \leq 0.78$ .

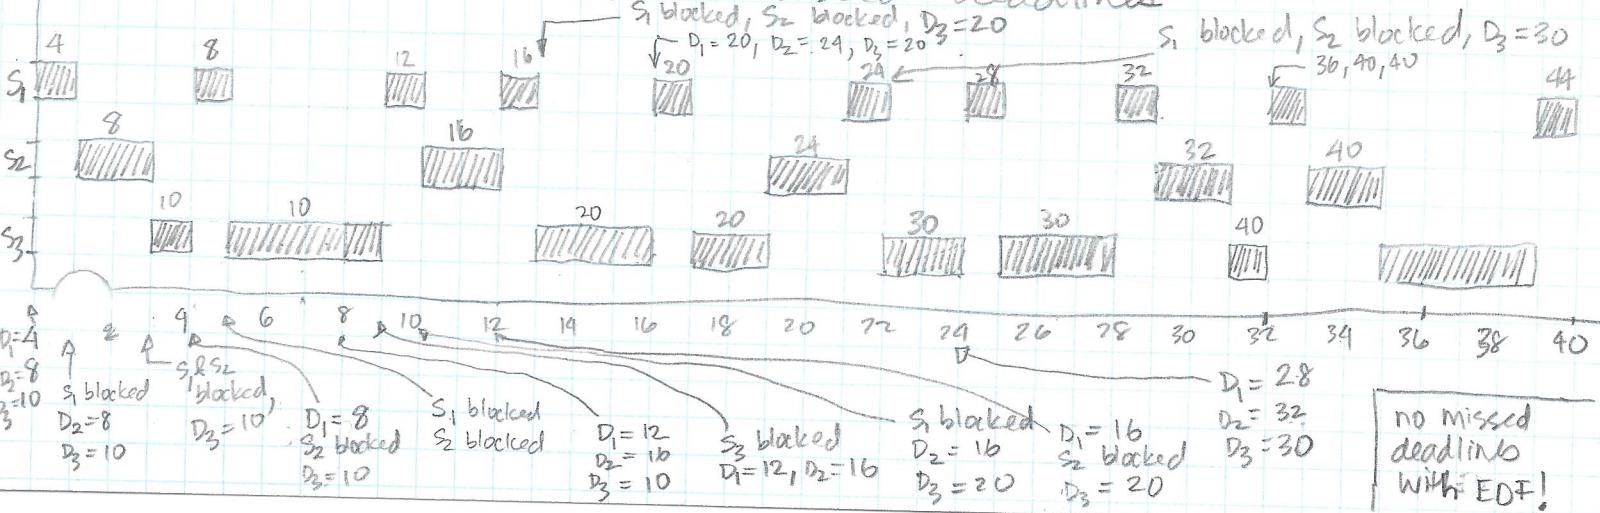
### EXERCISE (HOMEWORK/LAB) :

- in the previous example, set  $C_3 = 5$ . Is RM policy still optimal? How many missed deadlines are there over LCM time? Which task is most affected?
- apply EDF dynamic priority policy to this problem; are there missed deadlines?

### SOLUTIONS:



∴ there are three missed deadlines



FreeRTOS Overview

- FreeRTOS application form:

```
void Task1 (...) {
```

```
    while 1 {
```

```
        }
```

```
    return 1;
```

```
    }
```

```
    void Task2 (...) {
```

```
        while 1 {
```

```
            }
```

```
            return 1;
```

```
int main ( ) {
```

```
    prv.SetupHardware(); ← hardware initialization
```

```
    xTaskCreate (Task1, ...);
```

```
    xTaskCreate (Task2, ...);
```

```
    :
```

```
    vStartScheduler();
```

```
    return 1; ←
```

this  
statement  
is never  
reached

Tasks are  
created as  
infinite loops

main  
code  
in every  
FreeRTOS  
app looks  
something  
like  
this

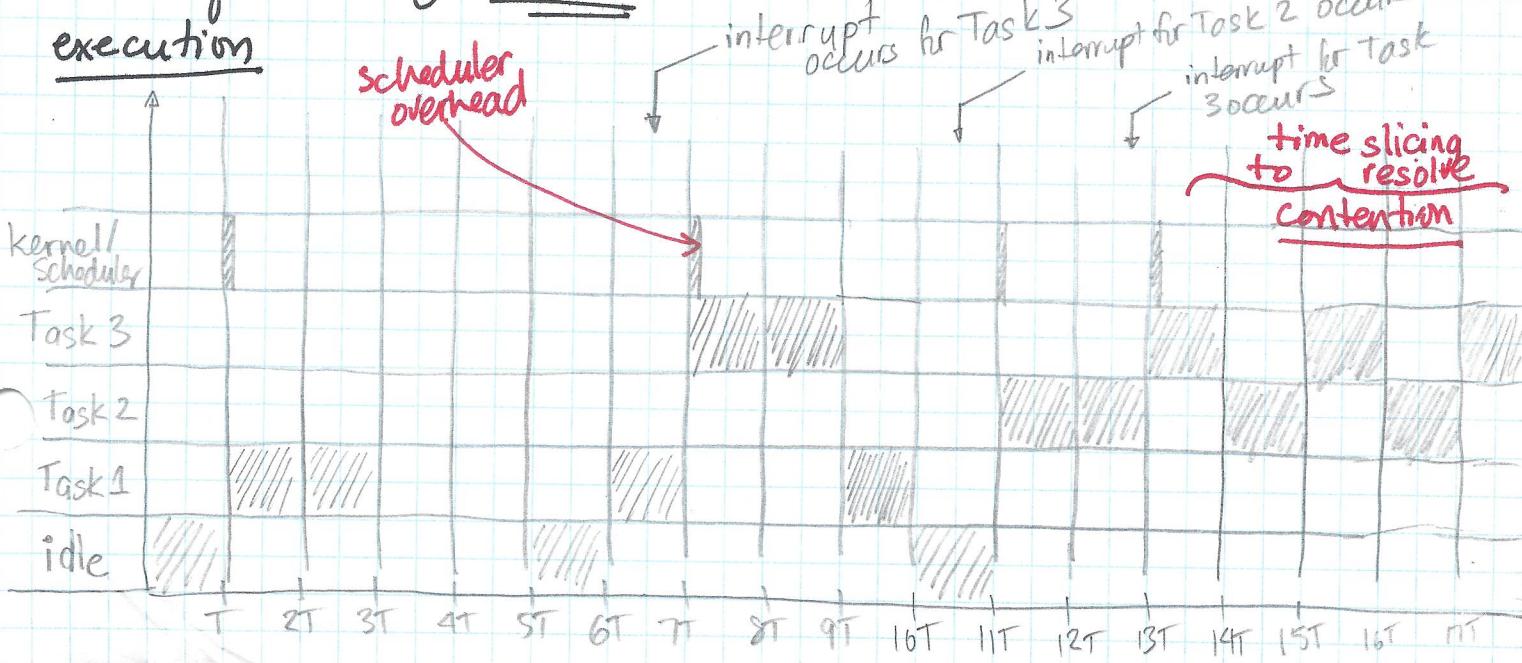
tasks are "setup"  
(not executed) here,  
including allocation of  
stack space on the  
heap and the setting  
of priorities

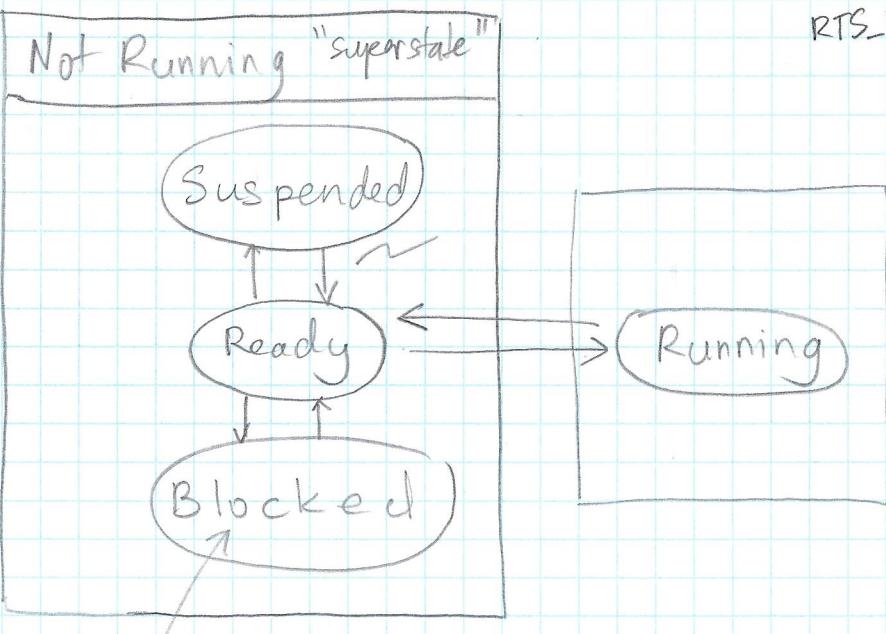
this is where the  
program hands control  
over to the Scheduler,  
which never relinquishes  
control

priorities

task	priority
:	:
Task 3	4
Task 2	4
Task 1	1
idle	0

- thus far we have only discussed setup → what happens before we hand control over to the scheduler
- now we can discuss execution → we assume: preemptive mode of operation (FreeRTOS can do cooperative as well), a single core, and a fixed set of priorities
- in addition, we assume a tick period which is many times (1000 to 1000 000 times the CPU clock period) — FreeRTOS has a default of 10 ms.





blocking occurs when the task is waiting for certain kinds of events, such as the expiration of a delay

→ why we use vTaskDelay(), a FreeRTOS function instead of another C-language delay implementation (which would keep the task in the running state).

- $\text{vTaskDelayUntil}(\cdot)$ :

## Task Delay ( $l-k$ )

function called

task calling vTaskDelay(.)

task calling `VTaskDelay()`  
transitions out of blocked state

$$0T_S \ 2T_S \ 3T_S \ 4T_S \ 5T_S \ 6T_S \quad \dots \ kT_S \quad \dots \ lT_S \quad (l+2)T_S$$

$$(k+1)T_S$$

1

$$(k+2)T_S$$

1

$$lT_S : (l+2)T_S$$

4

(l+1)TS

1

task which called  
v Task Delay Until (.) leaves  
blocked state

(assuming  $\&LastWakeTime = 0$ )

task calls vTaskDelayUntil(~~(f+2)s~~)

and ↑

~~Start~~ \$ x Last Wake Time

assumes a task is running periodically with a fixed frequency

ix Last Wake Time refers to the time at which the task last left the blocked state)